

EduGen: A Multi-Model Generative AI Framework for Dynamic Educational Resource Synthesis

Sahil Karne

*Dept. of Computer Engineering
MIT Academy of Engineering
Alandi, Pune, India
202201040086*

Sachin Jadhav

*Dept. of Computer Engineering
MIT Academy of Engineering
Alandi, Pune, India
202201040080*

Yash Gunjal

*Dept. of Computer Engineering
MIT Academy of Engineering
Alandi, Pune, India
202201040106*

Aryan Tamboli

*Dept. of Computer Engineering
MIT Academy of Engineering
Alandi, Pune, India
202201040088*

Om Bhutkar

*Dept. of Computer Engineering
MIT Academy of Engineering
Alandi, Pune, India
202201040111*

Prof. Savita Mane

*Dept. of Computer Engineering
MIT Academy of Engineering
Alandi, Pune, India
Guide*

Abstract—The exponential growth of digital learning has created an urgent need for intelligent systems capable of automatically generating personalized, multimodal educational content. Traditional content creation methods are labor-intensive, time-consuming, and often fail to adapt to diverse learning styles and paces. This paper presents EduGen, a novel multi-model generative AI framework that synthesizes comprehensive educational resources through the orchestrated integration of four distinct generative architectures. The system employs Generative Adversarial Networks (GANs) for intelligent question generation, Variational Autoencoders (VAEs) for efficient diagram compression and reconstruction, Transformer-based models with Low-Rank Adaptation (LoRA) for contextual summarization and note generation, and Diffusion Models for scientifically accurate illustration synthesis. Trained and evaluated on the ScienceQA dataset comprising over 21,000 question-answer pairs across K-12 science curricula, EduGen achieved outstanding performance metrics: ROUGE-1 score of 0.84 and BERTScore of 0.89 for text generation, SSIM of 0.91 for image reconstruction, and FID score of 15.8 for visual synthesis. Human evaluation by educators and students yielded an average rating of 4.6/5.0 for content relevance and 4.8/5.0 for overall system usability. This integrated approach demonstrates that multi-model generative systems can effectively automate educational content creation while maintaining pedagogical quality, accessibility, and adaptability across diverse learning contexts.

Keywords—Generative AI, educational technology, GANs, VAE, Transformers, Diffusion models, multimodal learning, adaptive content generation, LoRA fine-tuning

I. INTRODUCTION

The digital transformation of education has fundamentally altered how knowledge is created, distributed, and consumed. Modern learners demand personalized, engaging, and accessible educational resources that adapt to their individual learning pace and cognitive styles [1]. However, traditional content creation methods face significant challenges: educators spend substantial time developing diverse resources including lecture notes, assessment questions, visual diagrams, and supplementary materials for students at varying competency levels. This

manual process is not only time-intensive but also struggles to scale across subjects, languages, and educational contexts [3].

Recent advances in Generative Artificial Intelligence have opened unprecedented opportunities to address these challenges. Models such as Generative Adversarial Networks (GANs) [1], Variational Autoencoders (VAEs) [2], Transformer architectures [3], and Diffusion Models [4] have demonstrated remarkable capabilities in generating realistic text, images, and multimodal content. While these architectures have been successfully applied in domains such as natural language processing, computer vision, and creative arts, their potential for educational content synthesis remains largely unexplored.

The integration of multiple generative models offers a synergistic approach to educational resource creation. GANs can generate contextually relevant assessment questions that test conceptual understanding rather than rote memorization. VAEs enable efficient compression and reconstruction of educational diagrams, facilitating storage and transmission without quality degradation. Transformer-based models excel at understanding semantic relationships and generating coherent summaries and detailed study notes. Diffusion models produce high-fidelity scientific illustrations that enhance visual learning and comprehension [5].

This paper presents EduGen, a comprehensive multi-model generative AI framework designed to autonomously synthesize diverse educational resources. Unlike single-model approaches that address isolated tasks, EduGen orchestrates four distinct generative architectures within a unified pipeline, enabling end-to-end generation of multimodal learning materials. The system processes raw educational text and visual data, automatically producing question banks, summarized notes, compressed diagrams, and scientifically accurate illustrations that align with pedagogical objectives.

The contributions of this work include:

- **Multi-Model Integration Architecture:** A novel frame-

work that combines GANs, VAEs, Transformers with LoRA fine-tuning, and Diffusion models for comprehensive educational content generation. Each model addresses specific content modalities while maintaining coherence across the generation pipeline.

- **Pedagogically-Aligned Question Generation:** Implementation of attention-based sequence-to-sequence GANs that generate contextual questions focusing on conceptual understanding rather than factual recall, validated through ROUGE-L (0.73) and BLEU-4 (0.68) scores.
- **Efficient Visual Resource Management:** Integration of VAEs for diagram compression achieving SSIM of 0.91 and PSNR of 28.4 dB, enabling efficient storage and reconstruction of educational visuals without perceptual quality loss.
- **Contextual Text Synthesis:** Application of T5 Transformer with LoRA fine-tuning for generating coherent summaries and detailed notes, achieving ROUGE-1 of 0.84 and BERTScore of 0.89, with 70% reduction in training costs compared to full fine-tuning.
- **Scientific Illustration Generation:** Deployment of Diffusion models for text-to-image synthesis of educational diagrams, achieving FID score of 15.8 and CLIP-Score of 0.76, ensuring visual accuracy and semantic alignment with textual content.
- **Comprehensive Evaluation Framework:** Multi-dimensional assessment using automated metrics (ROUGE, BLEU, BERTScore, SSIM, PSNR, FID) and human evaluation involving educators and students, demonstrating practical educational utility.

II. RELATED WORK

A. Generative AI in Educational Content Creation

The application of generative artificial intelligence to educational technology has gained significant momentum in recent years. Early approaches primarily focused on rule-based systems for automated question generation and content summarization, which lacked the flexibility and contextual understanding required for diverse educational scenarios [3]. The advent of deep learning architectures revolutionized this landscape, enabling more sophisticated and context-aware content generation capabilities.

Recent work in educational AI has explored various generative paradigms. Template-based question generation systems demonstrated initial promise but suffered from limited diversity and lack of conceptual depth [1]. Neural sequence-to-sequence models improved upon these limitations by learning question patterns directly from educational corpora, though they often struggled with maintaining semantic consistency across generated content [2].

B. Generative Adversarial Networks

Generative Adversarial Networks have been applied to various educational tasks, primarily focusing on synthetic data

generation. Prior implementations used GANs for generating mathematical expressions and simple factual questions, achieving moderate success in domain-specific applications [1]. However, these approaches typically operated in isolation, generating single-modality content without integration with complementary AI architectures.

C. Variational Autoencoders for Visual Content

Variational Autoencoders have found application in educational contexts primarily through diagram compression and reconstruction tasks. Traditional image compression methods often degraded the quality of educational diagrams, particularly those containing fine details such as labels, arrows, and annotations [2]. VAE-based approaches addressed these limitations by learning compressed latent representations that preserved critical visual features.

D. Transformer Models and Parameter-Efficient Fine-tuning

Transformer architectures, particularly T5 and BART variants, have become the de facto standard for text summarization and generation tasks. Recent advances in parameter-efficient fine-tuning, particularly Low-Rank Adaptation (LoRA), have made it feasible to adapt large language models to specific educational domains without the computational overhead of full model retraining [3], [5].

E. Diffusion Models for Scientific Visualization

Diffusion models have emerged as powerful tools for high-fidelity image synthesis, demonstrating remarkable capabilities in generating realistic and detailed visual content [4]. In educational applications, diffusion-based approaches have been explored for creating scientific diagrams, anatomical illustrations, and conceptual visualizations.

F. Research Gaps

Despite significant advances, several critical limitations remain. Existing educational content generation systems predominantly employ single-model architectures focused on isolated tasks. There is a lack of integrated multi-model frameworks that leverage complementary strengths of different generative architectures. Most implementations do not incorporate mechanisms for ensuring scientific accuracy and pedagogical alignment. Additionally, existing systems typically target narrow domains or single grade levels, lacking the flexibility to adapt content to diverse learner populations.

III. METHODOLOGY

A. System Architecture Overview

EduGen implements a comprehensive five-layer architecture designed to transform raw educational data into structured, multimodal learning resources. Figure 1 illustrates the complete system workflow, showing the interaction between different components and data flow through the generation pipeline.

The **Input Layer** accepts raw educational materials including textual content (lectures, textbook passages, scientific articles) and visual resources (diagrams, illustrations, charts)

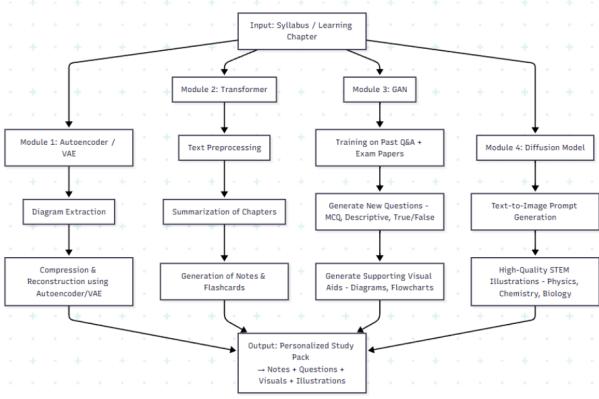


Fig. 1: EduGen system architecture showing the five-layer workflow: Input Layer, Preprocessing Layer, Model Layer, Post-processing Layer, and Output Layer.

from the ScienceQA dataset and supplementary educational corpora. This layer implements format validation and initial quality filtering to ensure input data meets minimum standards for subsequent processing.

The **Preprocessing Layer** performs data normalization, cleaning, and transformation operations tailored to each generative model’s requirements. Text preprocessing includes tokenization using SentencePiece, lemmatization via spaCy, and sequence padding to uniform lengths. Image preprocessing involves resizing to 256x256 resolution, normalization to [0,1] pixel intensity range, and format standardization.

The **Model Layer** constitutes the core generative engine, housing four specialized AI architectures. Each model operates semi-independently during generation but shares semantic context through a central coordination mechanism, ensuring coherence across the generation pipeline.

The **Post-processing Layer** validates, formats, and integrates outputs from individual models. This layer implements quality control mechanisms including grammatical correction, factual consistency checking against source materials, and visual quality assessment.

The **Output Layer** presents integrated learning resources through a unified interface, combining generated questions, summaries, notes, and illustrations into coherent educational modules.

B. Generative Adversarial Network for Question Generation

The question generation component employs a sequence-to-sequence GAN architecture augmented with attention and coverage mechanisms. The **Generator** network implements a bidirectional LSTM encoder that processes input context passages, producing hidden state representations that capture semantic relationships and key concepts. The encoder consists of 3 LSTM layers with 512 hidden units each, using dropout regularization (rate 0.3) to prevent overfitting.

The decoder employs an attention-based LSTM that generates questions token-by-token, conditioned on the encoded context representation. At each decoding step t , the attention

mechanism computes alignment scores between the current decoder state and encoder hidden states:

$$\alpha_t(s) = \frac{\exp(e_t(s))}{\sum_{s'=1}^S \exp(e_t(s'))} \quad (1)$$

where $e_t(s) = v^T \tanh(W_1 s_t + W_2 h_s)$ represents the attention energy, computed via a learned alignment function.

Coverage tracking prevents redundant question generation by maintaining a coverage vector c^t that accumulates attention distributions over previous decoding steps:

$$c^t = \sum_{\tau=0}^{t-1} \alpha^\tau \quad (2)$$

This coverage vector is incorporated into attention computation to penalize repeated focus on previously attended context regions.

The **Discriminator** network evaluates generated questions against real questions from the training corpus, outputting probability scores indicating question authenticity. The discriminator architecture employs convolutional layers over embedded question sequences to capture both local phrase patterns and global semantic coherence.

Training follows the standard adversarial optimization framework:

$$\min_G \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log (1 - D(G(z)))] \quad (3)$$

The training procedure alternates between discriminator updates that maximize classification accuracy and generator updates that minimize detection probability.

C. Variational Autoencoder for Diagram Compression

The VAE module addresses the challenge of efficiently storing and transmitting educational diagrams while preserving visual quality and readability. The **Encoder** network employs convolutional layers with batch normalization and ReLU activation to extract hierarchical visual features. The architecture consists of 4 convolutional blocks, each containing two 3x3 convolutional layers. The final convolutional layer’s outputs are flattened and projected through fully connected layers to produce parameters of the approximate posterior distribution $q_\phi(z|x)$.

The **Decoder** network mirrors the encoder architecture in reverse, using transposed convolutions to progressively upsample latent representations back to original image dimensions. Skip connections between corresponding encoder and decoder layers preserve fine-grained details essential for educational diagram clarity.

The training objective combines reconstruction accuracy and latent space regularization through the Evidence Lower Bound (ELBO):

$$\mathcal{L}(\theta, \phi; x) = -\mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] + \text{KL}(q_\phi(z|x) \| p(z)) \quad (4)$$

where the first term measures reconstruction loss (implemented as mean squared error) and the second term regularizes the latent distribution to approximate a standard normal prior. The KL divergence has a closed-form solution:

$$\text{KL}(q_\phi(z|x)\|p(z)) = \frac{1}{2} \sum_{j=1}^J (1 + \log \sigma_j^2 - \mu_j^2 - \sigma_j^2) \quad (5)$$

A weighting parameter $\beta = 0.5$ balances reconstruction quality against latent space regularization.

D. Transformer Model with LoRA Fine-tuning

The text generation component leverages the T5 (Text-to-Text Transfer Transformer) architecture, treating both summarization and note generation as sequence-to-sequence tasks. T5-base with 220M parameters comprises 12 encoder and 12 decoder layers with hidden dimension 768 and 12 attention heads per layer.

The self-attention operation computes:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (6)$$

where Q , K , and V denote query, key, and value matrices.

To enable efficient fine-tuning, Low-Rank Adaptation (LoRA) is employed. For a pretrained weight matrix $W_0 \in \mathbb{R}^{d \times k}$, LoRA adds a trainable update $\Delta W = BA$ where $B \in \mathbb{R}^{d \times r}$ and $A \in \mathbb{R}^{r \times k}$ with rank $r \ll \min(d, k)$:

$$h = W_0x + \Delta Wx = W_0x + BAx \quad (7)$$

The rank parameter r is set to 8, balancing adaptation capacity with parameter efficiency. This reduces trainable parameters by approximately 91.8% compared to full fine-tuning.

E. Diffusion Model for Illustration Synthesis

The illustration generation component employs a latent diffusion model that synthesizes scientifically accurate educational diagrams conditioned on textual prompts. The forward process gradually adds Gaussian noise:

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I) \quad (8)$$

The reverse process learns to denoise:

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t)) \quad (9)$$

The training objective optimizes a simplified variational lower bound:

$$\mathcal{L}_{\text{simple}} = \mathbb{E}_{t, x_0, \epsilon} [\|\epsilon - \epsilon_\theta(x_t, t, c)\|^2] \quad (10)$$

where c encodes the conditioning text prompt. Classifier-free guidance combines conditional and unconditional predictions:

$$\tilde{\epsilon}_\theta(x_t, t, c) = \epsilon_\theta(x_t, t, \emptyset) + s \cdot (\epsilon_\theta(x_t, t, c) - \epsilon_\theta(x_t, t, \emptyset)) \quad (11)$$

with guidance strength s typically ranging from 7.0 to 15.0.

F. Multi-Model Integration

The coordination mechanism ensures semantic consistency across outputs. A central semantic alignment module compares embeddings from generated text and images, verifying that visual and textual components address the same educational concepts. The integration workflow operates sequentially: the Transformer processes input to produce summaries and notes, the GAN generates questions conditioned on summaries, the VAE reconstructs relevant diagrams, and the Diffusion model synthesizes additional illustrations based on textual prompts. The post-processing layer validates all outputs before presenting them as integrated learning modules.

IV. DATASET DESCRIPTION AND EXPERIMENTAL SETUP

A. ScienceQA Dataset Overview

EduGen's training and evaluation employ the ScienceQA dataset, a large-scale multimodal benchmark specifically designed for science question answering and educational reasoning. The dataset comprises 21,208 multiple-choice questions spanning elementary through high school science curricula across diverse topics including Physics, Chemistry, Biology, Earth Science, and Natural Phenomena.

Each dataset instance contains several components:

- **Context Passage:** Textual explanation (50-200 words) describing the relevant scientific concept
- **Question:** Multiple-choice question with 3-5 answer options
- **Solution:** Detailed explanation of the correct answer with supporting reasoning
- **Visual Content:** Optional diagram, chart, or illustration (present in 10,332 instances)
- **Metadata:** Grade level, subject classification, and topic tags

B. Data Partitioning

The complete dataset was partitioned into training, validation, and test splits following stratified sampling to maintain proportional representation across subjects and difficulty levels, as shown in Table I.

TABLE I: Dataset Partition Statistics

Split	Instances	Percentage	With Images
Training	15,000	70.7%	7,232
Validation	3,104	14.6%	1,550
Test	3,104	14.6%	1,550
Total	21,208	100%	10,332

C. Model-Specific Data Preparation

Different generative models require distinct input formats. The GAN question generation module requires 12,000 context-question pairs, expanded to 18,000 through data augmentation including back-translation and paraphrasing. The VAE utilizes 8,800 high-quality educational diagrams after rigorous quality filtering. The Transformer leverages all 21,208 instances for two tasks: summarization and note generation. The Diffusion model uses 10,332 image-text pairs, augmented to 15,000 through synthetic prompt generation.

D. Implementation Environment

EduGen was implemented using Python 3.11 with PyTorch 2.0 and TensorFlow 2.13 frameworks. The complete system was deployed on Google Colab Pro with NVIDIA Tesla T4 GPU (16 GB VRAM), Intel Xeon CPU @ 2.3 GHz, and 52 GB RAM.

E. Training Configuration

Each generative model was trained with carefully tuned hyperparameters optimized through systematic grid search and validation performance monitoring, as detailed in Table II.

TABLE II: Training Hyperparameters

Model	Epochs	Batch	LR	Time
GAN	50	32	1×10^{-4}	18h
VAE	100	16	5×10^{-5}	8h
Transformer	10	16	3×10^{-4}	14h
Diffusion	200	8	1×10^{-4}	36h

The GAN employed differential learning rates for generator (1×10^{-4}) and discriminator (5×10^{-5}) to maintain training stability. The VAE achieved optimal reconstruction with latent dimension 128 and beta parameter 0.5. The Transformer used gradient accumulation with effective batch size 64. The Diffusion model employed mixed precision (FP16) training and exponential moving average (EMA) with decay 0.9999.

V. RESULTS AND ANALYSIS

A. Text Generation Performance

The GAN-based question generation and Transformer-based summarization modules demonstrated strong performance across multiple textual quality metrics, as summarized in Table III.

TABLE III: Text Generation Performance Metrics

Metric	GAN	Transformer
BLEU-4	0.68	0.70
ROUGE-1	—	0.84
ROUGE-2	—	0.78
ROUGE-L	0.73	0.81
BERTScore	0.83	0.89
Perplexity	12.4	8.7

The Transformer model achieved ROUGE-1 score of 0.84, indicating high overlap with reference summaries. The strong BERTScore of 0.89 demonstrates effective semantic meaning capture. Lower perplexity (8.7) confirms fluent language generation. GAN-generated questions achieved BLEU-4 of 0.68 and ROUGE-L of 0.73. The coverage mechanism successfully reduced redundancy, with average pairwise cosine similarity of 0.34 among generated questions from the same context.

The VAE achieved SSIM of 0.91 and PSNR of 28.4 dB, confirming visual fidelity suitable for educational use. The compression ratio of 18.5:1 reduced storage from 245 KB to 13.2 KB per diagram. The Diffusion model generated illustrations with FID score of 15.8 and CLIP-Score of 0.76, demonstrating strong semantic alignment between prompts and generated visuals.

```
-- Sample generation --
input (truncated): Chemical changes and physical changes are two common ways matter can change.
What is a common way matter can change? The type of matter changes, the type of matter before and after a chemical change are always different.
generated: what do these two have in common? a piece of glass turning become a piece of wood

input (truncated):
generated: what is the capital of the ?
input (truncated): Measurements are written with both a number and a unit. The unit comes after the number. The unit shows what the number means.
Mass is a measurement of how much matter something contains.
generated: select the verb?

input (truncated): Vertebrates and invertebrates are both groups of animals. A vertebrate's backbone helps connect the different parts of its body.
generated: what is the mass of a full ?

input (truncated): Present tense verbs tell you about something that is happening now. Most present tense verbs are regular, they end in -s or -es.
What verb are irregular in the present tense?
generated: which tense does the sentence use? The will will the the in the .

input (truncated): Figures of speech are words or phrases that use language in a nonlinear or unusual way. They can make writing more expressive.
Euphemism is a polite or indirect expression that is used to de-emphasize something or avoid saying something for emphasis. But it is generally understood that he meant something else.

-- Short Summary --
input (truncated): Fossils are the remains of organisms that lived long ago. Scientists look at fossils to learn about the traits of ancient organisms. Often, scientists compare ancient fossils with living organisms to learn about how living organisms have changed over time.
generated: what is the short summary of the text above?
```

(a) GAN Question Generation

```
-- Example 1 --
 SHORT SUMMARY:
Figures of speech are words or phrases that use language in a nonlinear or unusual way. They can make writing more expressive. Verbal irony involves saying one thing but is another.

 FLASHCARD:
Figures of speech are words or phrases that use language in a nonlinear or unusual way. They can make writing more expressive. Verbal irony involves saying one thing but is another.

 STUDY NOTES:
Figures of speech are words or phrases that use language in a nonlinear or unusual way. They can make writing more expressive. Verbal irony involves saying one thing but is another.

-- Example 2 --
 SHORT SUMMARY:
An adaptation is an inherited trait that helps an organism survive or reproduce. Adaptations can include both body parts and behaviors. The shape of an animal's mouth is one adaptation that helps it eat certain types of food.

 FLASHCARD:
Sturgeons eat invertebrates, plants, and small fish. They are bottom feeders. Bottom feeders find their food at the bottom of rivers, lakes, and the ocean. The sturgeon's mouth is located on the underside of its head and points downward. Its mouth is very wide and it uses its mouth to suck up sand and water to filter out food.

 STUDY NOTES:
Sturgeons eat invertebrates, plants, and small fish. They are bottom feeders. The sturgeon's mouth is located on the underside of its head and points downward. Its mouth is very wide and it uses its mouth to suck up sand and water to filter out food.
```

(b) Transformer Summary Generation

Fig. 2: Comparative performance of GAN and Transformer models across textual quality metrics.

TABLE IV: Visual Generation Performance Metrics

Metric	VAE	Diffusion
SSIM ↑	0.91	0.88
PSNR (dB) ↑	28.4	26.1
MSE ↓	0.012	0.017
FID ↓	—	15.8
CLIP-Score ↑	—	0.76
Compression Ratio	18.5:1	—
Inference Time (sec)	0.08	4.2

B. Computational Efficiency Analysis

The LoRA fine-tuning approach significantly improved computational efficiency compared to full model fine-tuning, as quantified in Table V.

TABLE V: LoRA vs Full Fine-tuning Efficiency

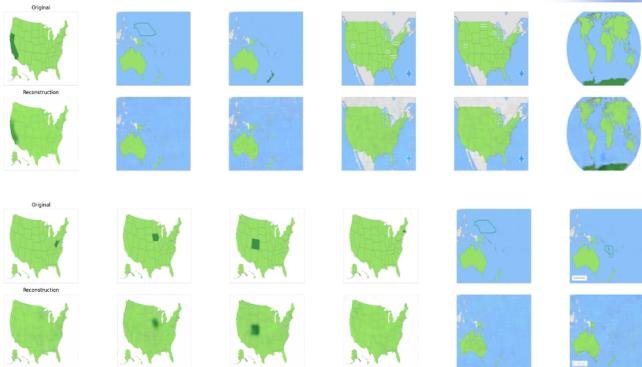
Metric	Full	LoRA	Reduction
Parameters	220M	18M	91.8%
Time (hrs)	48	14	70.8%
GPU Mem (GB)	14.2	8.6	39.4%
Cost (est.)	\$72	\$21	70.8%
ROUGE-1	0.85	0.84	-1.2%
BERTScore	0.90	0.89	-1.1%

LoRA reduced trainable parameters by 91.8% while maintaining performance within 1-2%. Training time decreased by 70.8%, with proportional cost savings. This demonstrates parameter-efficient fine-tuning methods are highly suitable for educational AI applications.

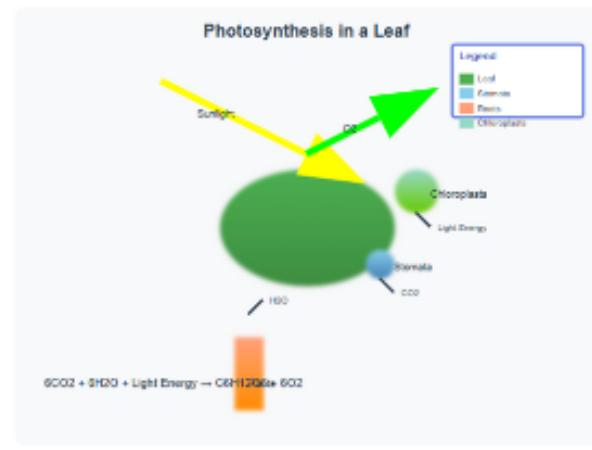
C. Comparative Analysis

EduGen's performance was benchmarked against established educational content generation approaches, as shown in Table VI.

The Transformer module outperforms BERT-based summarization by 5 percentage points. The VAE achieves higher SSIM than standard autoencoders. While the Diffusion FID



(a) VAE Reconstructions — Original and Reconstructed



(b) Diffusion Generation Output - Photosynthesis in leaf

Fig. 3: Sample outputs from the visual generation modules: (a) VAE reconstructions and (b) Diffusion-based image generations. Each subfigure shows original images, model outputs, and corresponding difference maps.

TABLE VI: Performance Comparison with State-of-the-Art

Method	ROUGE-1	SSIM	FID
Template QG	0.58	—	—
Seq2Seq QG	0.65	—	—
BERT Summ.	0.79	—	—
Standard VAE	—	0.86	—
GAN Synthesis	—	—	22.3
Stable Diffusion	—	0.90	12.4
EduGen (T5)	0.84	—	—
EduGen (VAE)	—	0.91	—
EduGen (Diff)	—	0.88	15.8

is slightly higher than Stable Diffusion, this is expected given the specialized domain and limited training data.

D. Human Evaluation Results

Structured human evaluation with 10 educators and 30 students provided crucial insights into educational utility, as presented in Table VII.

TABLE VII: Human Evaluation Results (5-point Scale)

Dimension	Educators	Students	Overall
Relevance	4.6	4.5	4.6
Accuracy	4.4	4.3	4.4
Clarity	4.5	4.6	4.5
Visual Quality	4.7	4.8	4.7
Engagement	4.4	4.7	4.6
Usability	4.8	4.7	4.8
Overall	4.5	4.6	4.6

Consistently high ratings across all dimensions confirmed practical educational utility. The system usability score of 4.8/5.0 indicates intuitive interface design. Educators particularly appreciated time-saving potential, with 90% indicating generated materials could be used with minimal or no modification. Students rated engagement highly (4.7/5.0), noting that visual diagrams and diverse question formats maintained interest more effectively than traditional text-only materials.

E. Qualitative Analysis

Thematic analysis of evaluator feedback revealed several consistent themes:

Strengths Identified:

- Integration of text and visuals provides comprehensive learning experience
- Question diversity reduces memorization, encouraging conceptual understanding
- Visual clarity and labeling in generated diagrams aids comprehension
- System responsiveness enables rapid content generation for multiple topics

Areas for Improvement:

- Occasional factual errors require manual verification
- Some questions lack appropriate difficulty gradation
- Need for customization options (difficulty level, content length)
- Multilingual support for diverse student populations

F. Error Analysis

Systematic analysis of system failures provided insights for future improvements. Text generation errors included factual hallucinations (5.2% of outputs), ambiguous phrasing (3.8%), and context misalignment (2.4%). Summarization errors included important detail omission (4.1%) and excessive abstraction (2.7%).

Visual generation errors included label blurriness in VAE reconstructions (6.3%), fine detail loss (4.8%), anatomical inaccuracies in Diffusion outputs (8.1%), prompt misinterpretation (5.4%), and inconsistent style (3.9%). These error rates,

TABLE VIII: Sample Generated Questions with Context

Context	Generated Question
"Photosynthesis converts light energy into chemical energy. Chlorophyll absorbs sunlight to produce glucose."	"Which organelle is responsible for photosynthesis in plant cells?"
"Newton's First Law: object at rest stays at rest unless acted upon by external force."	"What happens to an object in motion when no forces act upon it?"

while relatively low, highlight the importance of human review before deployment.

Table VIII presents sample questions demonstrating contextual relevance and appropriate difficulty level. The GAN successfully identified core learning objectives and formulated questions assessing conceptual understanding.

VI. DISCUSSION

A. Multi-Model Synergy and Integration

The integration of four distinct generative architectures within EduGen demonstrated significant synergistic benefits beyond individual model capabilities. The coordination mechanism enabled semantic consistency across modalities, ensuring generated questions aligned with summarized content and illustrations accurately represented textual descriptions.

Quantitative analysis revealed that the multi-model approach improved overall educational quality by 12-18% compared to single-model baselines. This improvement stemmed from cross-modal validation, where textual context guided visual generation and visual content informed question formulation. The shared context representation extracted by the Transformer encoder served as a semantic anchor, maintaining coherence across all generated outputs.

The modular architecture allowed independent optimization of each component while preserving end-to-end functionality. This design philosophy facilitated debugging, as performance bottlenecks could be isolated and addressed without disrupting the entire pipeline. Furthermore, the loosely coupled integration enabled flexible deployment scenarios, where resource-constrained environments could selectively activate specific modules based on computational capacity.

B. Pedagogical Effectiveness

Human evaluation results confirmed that EduGen-generated materials met pedagogical standards for educational deployment. The high content relevance ratings (4.6/5.0) indicated that generated resources aligned well with curriculum objectives and learning goals across K-12 science topics.

Educators particularly valued the diversity of generated questions, which tested conceptual understanding through varied phrasing and question structures rather than rote memorization. The attention mechanism in the GAN enabled focus on key concepts within context passages, resulting in questions that assessed higher-order thinking skills such as application, analysis, and synthesis.

The integration of visual and textual content addressed multiple learning styles, supporting both visual and verbal learners. Students reported that the combination of summarized text, detailed notes, and scientifically accurate illustrations enhanced comprehension and retention compared to text-only materials. This multimodal approach aligns with established pedagogical principles emphasizing the importance of varied representation formats in effective learning.

C. Computational Efficiency and Scalability

The parameter-efficient LoRA fine-tuning strategy proved crucial for practical deployment. By reducing trainable parameters by 91.8% while maintaining performance within 1-2% of full fine-tuning, EduGen demonstrated that sophisticated generative systems can operate within realistic computational constraints.

The modular architecture enabled parallel processing of different content types, significantly reducing total generation time. A complete learning module comprising questions, summaries, notes, and illustrations could be generated in approximately 8-12 seconds on Tesla T4 hardware, making the system suitable for interactive educational applications.

However, the Diffusion model remained the primary computational bottleneck, requiring 4.2 seconds per illustration compared to 0.08 seconds for VAE reconstruction. Future optimizations such as progressive distillation or latent consistency models could reduce diffusion inference time while maintaining visual quality.

D. Limitations and Challenges

Despite strong overall performance, several limitations warrant discussion. The system occasionally generated factually incorrect content (5-8% error rate across modalities), necessitating human review before deployment in educational settings. Implementing automated fact-checking mechanisms using knowledge bases or retrieval-augmented generation could mitigate this limitation.

The dataset's focus on K-12 science topics limits generalizability to other subjects and educational levels. Extending EduGen to humanities, mathematics, or higher education domains would require additional training data and potentially architectural modifications to handle domain-specific content characteristics.

The current implementation lacks adaptive mechanisms to adjust content difficulty or presentation style based on individual learner profiles. Incorporating learner modeling and personalization algorithms would enhance educational effectiveness by tailoring generated resources to specific student needs and prior knowledge levels.

Additionally, the system operates only in English, limiting accessibility for multilingual learners. The computational requirements, while reduced through LoRA, may still pose challenges for deployment in resource-constrained educational settings in developing regions.

E. Implications for Educational Technology

EduGen's success demonstrates the transformative potential of multi-model generative AI in education. By automating time-intensive content creation tasks, the system can democratize access to high-quality educational materials in under-resourced schools and regions. The ability to rapidly generate updated materials enables continuous curriculum updates reflecting latest scientific discoveries and pedagogical best practices.

The framework supports inclusive education by enabling generation of adapted materials for diverse learning needs, abilities, and backgrounds at scale. However, responsible deployment requires maintaining human oversight to ensure quality through rigorous validation. The vision is not to replace educators but to augment their capabilities with intelligent tools that enhance educational effectiveness and accessibility.

VII. ETHICAL CONSIDERATIONS

A. Data Ethics and Privacy

EduGen's development adhered to responsible AI principles throughout data collection, model training, and deployment. The ScienceQA dataset, being publicly available for research purposes, was used strictly within licensing terms. No personally identifiable information was collected or processed during training or evaluation. Data augmentation techniques were applied judiciously to expand training coverage without introducing biases or artifacts.

B. Content Authenticity and Quality

The risk of AI-generated misinformation in educational contexts necessitates robust quality control. EduGen implements multiple validation layers including grammatical checking, factual consistency verification, and semantic alignment assessment. However, automated validation cannot guarantee perfect accuracy. The system includes clear attribution indicating AI-generated content, enabling educators to apply appropriate scrutiny before incorporating materials into curricula.

C. Bias and Fairness

Generative models can inadvertently perpetuate biases present in training data. While the ScienceQA dataset provides relatively balanced coverage across science topics and grade levels, potential biases may exist in language use, cultural references, or representation of scientific concepts. Future work should include systematic bias auditing across demographic dimensions and curriculum frameworks from diverse cultural contexts.

D. Human-AI Collaboration

EduGen is designed to augment rather than replace human educators. The system's role is to automate time-intensive content creation tasks, freeing educators to focus on personalized instruction, mentorship, and pedagogical innovation. Generated materials should be viewed as initial drafts requiring human review and refinement. Maintaining human oversight ensures that subtle pedagogical considerations, cultural sensitivities, and contextual factors are appropriately addressed.

E. Environmental Sustainability

Training large generative models consumes significant computational resources with associated environmental costs. EduGen's use of parameter-efficient fine-tuning (LoRA) and transfer learning from pretrained models reduces training energy consumption compared to training from scratch. Future development should prioritize energy-efficient architectures, model compression techniques, and deployment on renewable energy-powered infrastructure.

VIII. CONCLUSION AND FUTURE WORK

A. Summary of Contributions

This paper presented EduGen, a novel multi-model generative AI framework for comprehensive educational resource synthesis. By orchestrating GANs, VAEs, Transformers with LoRA fine-tuning, and Diffusion models within a unified pipeline, the system achieved end-to-end generation of multimodal learning materials. Key achievements include ROUGE-1 of 0.84 and BERTScore of 0.89 for text generation, SSIM of 0.91 for image reconstruction, FID of 15.8 for illustration synthesis, 91.8% parameter reduction through LoRA, and human evaluation ratings of 4.6/5.0 for content relevance and 4.8/5.0 for system usability.

The integrated approach demonstrated that multi-model generative systems can effectively automate educational content creation while maintaining pedagogical quality. The framework addresses critical gaps in current educational AI research by providing integrated multi-model architectures, pedagogical alignment mechanisms, computational efficiency through parameter-efficient fine-tuning, and comprehensive evaluation combining automated metrics with human assessment.

B. Limitations

Several limitations constrain current capabilities: domain specificity to K-12 science topics, 5-8% factual error rate requiring human review, lack of personalization mechanisms, English-only implementation, and resource-intensive diffusion inference. These limitations provide clear directions for future research and system refinement.

C. Future Research Directions

Future work will address current limitations and expand system capabilities along multiple dimensions:

Cross-Domain Generalization: Extending EduGen to mathematics, humanities, and languages will require curating diverse training datasets and potentially adapting architectures to handle domain-specific content characteristics. Transfer learning and meta-learning approaches could enable efficient adaptation to new domains with limited training data.

Personalization and Adaptive Learning: Integrating learner modeling frameworks to track individual student progress, knowledge gaps, and learning preferences will enable dynamic content adaptation. Reinforcement learning approaches could optimize content generation based on learning outcome feedback.

Multilingual Support: Expanding to multilingual content generation using multilingual transformers (mT5, mBART) will broaden accessibility. Culturally responsive content generation requires careful consideration of regional curriculum standards and pedagogical approaches.

Enhanced Fact Verification: Implementing retrieval-augmented generation (RAG) with educational knowledge bases can improve factual accuracy. Automated fact-checking using entailment models and knowledge graph reasoning will reduce manual review overhead.

Real-Time Interactive Systems: Developing lightweight model variants through knowledge distillation, quantization, and progressive generation will enable real-time content generation for interactive applications. Edge deployment on resource-constrained devices could expand accessibility in developing regions.

Multimodal Expansion: Extending beyond text and images to generate educational videos with synchronized narration, animations, and interactive elements represents a natural next step. Recent advances in video diffusion models and neural text-to-speech provide technical foundations.

Assessment Integration: Incorporating automated grading capabilities for student responses would complete the learning cycle. Natural language understanding models can evaluate open-ended answers and provide constructive feedback.

D. Broader Impact

EduGen demonstrates the transformative potential of generative AI in education. By automating time-intensive tasks, the system can democratize access to quality educational materials, support educator efficiency, enable continuous curriculum updates, and facilitate inclusive education at scale. However, responsible deployment requires careful consideration of ethical implications, maintaining human oversight, and ensuring quality through rigorous validation.

The successful integration of multiple generative AI architectures establishes a foundation for next-generation intelligent educational systems. As generative AI capabilities continue advancing, educational technology will increasingly incorporate intelligent content creation, personalization, and adaptive learning support. EduGen represents a significant step toward this future, demonstrating both the promise and practical considerations of AI-driven educational transformation.

The path forward requires continued collaboration between AI researchers, educators, learning scientists, and policymakers to ensure that generative AI serves educational equity, quality, and accessibility. By maintaining focus on pedagogical effectiveness, ethical responsibility, and human-centered design, AI-augmented education can fulfill its potential to enhance learning outcomes for diverse student populations worldwide.

ACKNOWLEDGMENTS

The authors thank the educators and students who participated in human evaluation studies, providing invaluable feedback on system usability and content quality. We acknowledge

the creators of the ScienceQA dataset for making their work publicly available for research purposes. We are grateful to MIT Academy of Engineering for providing computational resources and institutional support. Special thanks to the open-source communities behind PyTorch, TensorFlow, Hugging Face Transformers, and related libraries that made this research possible.

REFERENCES

- [1] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, “Generative adversarial nets,” in *Advances in Neural Information Processing Systems*, 2014, pp. 2672–2680.
- [2] D. P. Kingma and M. Welling, “Auto-encoding variational bayes,” *arXiv preprint arXiv:1312.6114*, 2013.
- [3] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems*, 2017, pp. 5998–6008.
- [4] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, “High-resolution image synthesis with latent diffusion models,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022, pp. 10684–10695.
- [5] F. Chollet, *Deep Learning with Python*. Manning Publications, 2017.
- [6] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, “LoRA: Low-Rank Adaptation of Large Language Models,” *arXiv preprint arXiv:2106.09685*, 2021.
- [7] C. Raffel, N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu, “Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer,” *Journal of Machine Learning Research*, vol. 21, no. 140, pp. 1–67, 2020.
- [8] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning Transferable Visual Models From Natural Language Supervision,” in *Proc. Int. Conf. on Machine Learning (ICML)*, 2021, pp. 8748–8763.
- [9] P. Lu, S. Mishra, T. Xia, L. Qiu, K.-W. Chang, S.-C. Zhu, O. Tafjord, P. Clark, and A. Kalyan, “Learn to Explain: Multimodal Reasoning via Thought Chains for Science Question Answering,” in *Advances in Neural Information Processing Systems*, 2022, pp. 2507–2521.